

DNS-Based Traffic Steering for IPv6-Enabled Edge Microservices

Extended Project Description

Joar Heimonen
contact@joar.me

May 22, 2026

Contents

1	Introduction	2
2	State of the Art in Traffic Steering and Proxy Scalability	2
2.1	The End-to-End Principle and the Role of NAT	2
2.2	Layer 7 Proxies and Ingress Controllers	3
2.3	Service Mesh Architectures	3
2.4	Kernel and Hardware Acceleration	3
2.5	DNS-Based Traffic Steering	4
2.6	Edge and Orchestration Context	5
3	Literature Search Methodology	5
4	Problem Description	6
5	Discussion	6
6	Proposed Approach and Methodology	6
7	Work Plan and Milestones	6
8	Conclusion	7

1 Introduction

The modern cloud computing stack relies heavily on Layer 7 proxies and service meshes to manage traffic between services however this comes at a cost. These layer 7 proxies introduce single points of failure in otherwise distributed systems, whilst adding unnecessary complexity and reducing the effectiveness of layer 3 routing. Service meshes add sidecar containers to every pod. This overhead scales with the amount of microservices you have.

This architectural pattern did not emerge by design. IPv4 exhaustion forced the development of technologies like NAT (Network Address Translation) and shared address spaces. With the introduction of NAT end-to-end accessibility was broken making proxies a structural necessity.

With a rapid increase in IPv6 adoption it is time to leave the technical limitations of the past in the past and build our services with IPv6s large address space in mind. Under RIPE-738 [1], a Local internet registry receives a minimum allocation size of CIDR /32 up to /29 without needing to justify it. A /29 allocation alone contains around 147 billion times more addresses than the whole IPv4 address space.

This thesis investigates whether DNS-based traffic steering, combined with direct IPv6 addressing of microservices can serve as a viable alternative to traditional layer 7 ingress proxies and service meshes. Rather than routing traffic through a centralized proxy each pod is assigned a globally routable address and made directly reachable by clients. Traffic steering decisions are delegated to a purpose built DNS server that dynamically manages record based on pod health and metrics.

The remainder of this document surveys the state of the art in traffic steering and proxy scalability, identifies the gap that motivates this research and outlines the proposed approach, experimental methodology and work plan.

2 State of the Art in Traffic Steering and Proxy Scalability

2.1 The End-to-End Principle and the Role of NAT

Saltzer, Reed and Clark established in 1984 that application specific functions should be implemented at the endpoints of a communication system rather than within the network itself [2]. A network that simply moves packets without enforcing control points is inherently more resilient. There are no forced bottlenecks, no single points of failure introduced by the infrastructure itself. Packets are free to be routed as the network sees fit.

The widespread adoption of NAT under IPv4 fundamentally broke this principle. By placing address translation logic inside the network, NAT made end-to-end addressability impossible. Applications could no longer communicate directly over the internet. They required proxies, STUN servers and other middleboxes to facilitate communication compensating for the loss of globally unique addressing.

IPv6 restores the conditions under which end-to-end addressability principle can be applied. With a globally unique address assigned to every endpoint there is no longer a structural requirement for address translation or centralized proxies. The network can return to its intended role, moving packets between endpoints that are directly reachable by design.

2.2 Layer 7 Proxies and Ingress Controllers

A reverse proxy sits between the client and the backend services as a mandatory intermediary. This intermediary is responsible for TLS termination, load balancing, health checking and routing. These are functions that every service in the system depends on, making the proxy both a critical component and an inherent bottleneck.

HAProxy, NGINX, Envoy and Traefik represent the dominant L7 proxy implementations. Despite their differences in design and feature set, they share a common architectural assumption. All traffic must pass through a userspace process that terminates connections, inspects application-layer content, and forwards requests to backend services. Wang, Shou, Qian and Liu quantify the cost of this assumption, sidecar-based proxies reduce throughput by up to 55.9% and more than double latency even under light load, with only 20% of that overhead attributable to actual load balancing logic [3].

The centralization is the problem. All traffic passes through one process that becomes a bottleneck and a single point of failure in an otherwise distributed system. The more services present the worse it gets.

2.3 Service Mesh Architectures

Service meshes are a backbone that is supposed to allow services to arbitrarily talk to each other. Most modern service meshes can be seen as consisting of two planes. A data plane and a control plane. The data plane consists of sidecar proxies injected into every pod, whilst the control plane distributes configurations, certificates and routing rules to the sidecar proxies whilst maintaining a global graph of the mesh.

When pod A calls pod B the traffic flows from pod A to pod A's sidecar proxy, through the network and then to pod B's sidecar proxy. The sidecar proxy then relays the call to the pod. On top of this the control plane is constantly rotating certificates, changing routing rules, updating configurations and managing service discoverability. It becomes its own distributed system just for passing messages between nodes of another distributed system.

Li, Lu, Lin, Wu, Zhang and Yan measure the performance of this architecture directly. Deploying a service mesh more than doubles communication latency, decreases throughput by over 60%, and consumes over 70% of CPU resources [4]. The large overhead is not incidental, it is a consequence of routing every inter-service call through multiple user-space processes, each crossing the kernel boundary multiple times per request. This results in an alarming amount of context switches for a system that is supposed to be efficient.

2.4 Kernel and Hardware Acceleration

A recent development in the field is the moving of L7 logic into the kernel. Rather than eliminating the proxy, these approaches attempt to reduce its overhead by removing the large overhead caused by crossing the kernel user-space boundary multiple times per request.

Wang, Shou, Qian and Liu propose XLB, an in kernel L7 load balancer implemented using eBPF (Extended Berkley Packet Filtering). By intercepting traffic in the socket layer rather than routing it through a separate user-space process XLB reduces the cross process communication overhead that makes traditional sidecar proxies expensive. The result is up to 1.5x higher

throughput and 60% lower latency [3]. However, this approach is currently not viable as it requires re-implementing user-space protocols in the kernel. XLB only supports unencrypted HTTP 1.1 and support for newer protocols would require the implementation of a full in kernel layer 7 network stack, including support for websockets, webtransports, HTTP 3.0 (quic support) and TLS support.

Miziński and Przyłucki confirm the broader trend, finding that Cilium’s eBPF based networking reduces CPU and memory consumption significantly under load compared to traditional iptables based approaches. however, this comes at a slight cost to raw throughput [5].

Li, Lu, Lin, Wu, Zhang and Yan propose FlatProxy, which migrates the service mesh proxy onto a Data Processing Unit (DPU). A DPU is a dedicated hardware component, in FlatProxy’s case it is a Field Programmable Gate Array (FPGA). The FlatProxy reduces overhead by running the proxy on the DPU rather than the CPU. However, to avoid implementing L7 protocols on an FPGA they have devised a simpler solution. The DPU routes traffic that it has not seen before to a user space process on FlatProxy’s CPU. The user-space process then creates routing rules that are applied to all future proxies by the DPU. FlatProxy eliminates resource contention between proxy and application, achieving a 90% reduction in latency and 4x improvement in throughput compared to Envoy [4]. However, DPU-based approaches introduce significant hardware requirements that make them unsuitable for commodity or edge deployments.

Both hardware accelerated proxies and in kernel proxies share a fundamental assumption with the systems they are trying to improve upon: traffic must still pass through a dedicated proxy. The overhead is reduced but, the centralization remains the same.

2.5 DNS-Based Traffic Steering

The Domain Name System (DNS) is part of the backbone of the internet, It is what allows our services to remain static as the address space evolves. Rather than routing traffic through a proxy DNS-based traffic steering directs clients to the appropriate backend before the connection is even established. The steering decisions happens at or before resolution time instead of at request time.

Hawley demonstrated DNS-based geographical traffic steering in 2009 with GeoDNS, deployed at kernel.org to distribute Linux kernel downloads across mirrors in multiple countries [6]. The approach is protocol agnostic and completely transparent to the servers and clients. Notably, DNS-based steering was not considered a fringe idea at the time. Wikipedia, Mozilla and the Apache foundation all operated GeoDNS infrastructure at the time. The concept predates all modern cloud paradigms.

DNS is hierarchical. Responsibility for a zone is delegated from parent to child through a chain of authority. At the top of this chain sits the root zone, delegating to top-level domains, which delegate to authoritative nameservers for individual domains. Each zone has a Start Of Authority (SOA) record that identifies the primary nameserver responsible for that zone and defines the default TTL value for its records.

The canonical limitation of DNS-based steering is TTL caching. Moura, Heidemann and Schmidt demonstrate that recursive resolvers frequently cache records beyond their stated TTL, and that the effective cache lifetime of a record can far exceed what the operator intended [7]. Shaikh, Tewari and Agrawal show that reducing TTL to improve steering responsiveness increases name-server query load by orders of magnitude [8]. This creates a fundamental tension between steering

freshness and infrastructure cost.

This may be mitigated through proper delegation of authority. By delegating a dedicated subzone to its own authoritative nameserver and ensuring that recursive resolvers are only served NS records for that zone, clients are forced to contact the authoritative nameserver directly. The stub resolver on the client machine, upon receiving an NS record, bypasses the recursive resolver’s cache entirely and queries the authoritative server. This means steering decisions are always fresh, regardless of what recursive resolvers choose to cache, because the client never receives a cacheable A or AAAA record from the recursive resolver. However, modern DNS stub resolvers and recursive servers often only loosely adhere to DNS standards, making the practical effects of fringe delegation mechanisms difficult to predict.

No one has applied this at the per-pod level in a microservice environment.

2.6 Edge and Orchestration Context

3 Literature Search Methodology

The literature search was conducted using Google Scholar and Google Scholar Labs. Search terms were initially generated with the assistance of a large language model based on the research topic, and subsequently refined through iterative searching. The following keyword queries were used:

- *DNS TTL load balancing consistency tradeoff*
- *IPv6 end-to-end addressing NAT elimination*
- *eBPF Kubernetes networking performance*
- *service mesh overhead scalability microservices*
- *DNS-based server selection load balancing*

Papers were selected based on relevance to the research problem, publication venue quality, and citation count. The search was supplemented by backward snowballing from the reference lists of key papers, through which several foundational works were identified. Papers published in predatory or low-quality journals were excluded.

The resulting set of nine papers is summarized in Table 1.

#	Authors	Year	Key Claim
1	Wang, Shou, Qian, Liu	2026	In-kernel eBPF L7 load balancer, 1.5x throughput over Istio [3]
2	Li, Lu, Lin, Wu, Zhang, Yan	2023	DPU-centric service mesh, 90% latency reduction vs Envoy [4]
3	Hawley	2009	GeoDNS geographic traffic steering at Internet scale [6]
4	Jeffery, Howard, Mortier	2021	etcd strong consistency bottleneck in edge Kubernetes [9]
5	Miziński, Przyłucki	2025	Cilium eBPF vs iptables: lower CPU/memory under load [5]
6	Moura, Heidemann, Schmidt	2019	DNS TTL tradeoffs between caching efficiency and routing flexibility [7]
7	Shaikh, Tewari, Agrawal	2001	Reducing DNS TTL increases name-server load by orders of magnitude [8]
8	Chatterjee, Tari, Zomaya	2005	Adaptive TTL approach reduces client-perceived latency by 16% [10]
9	Saltzer, Reed, Clark	1984	End-to-end principle: functions belong at endpoints, not in the network [2]

Table 1: Papers included in the literature review

4 Problem Description

5 Discussion

6 Proposed Approach and Methodology

7 Work Plan and Milestones

Spring 2026

- Systematization of knowledge and literature study
- Write and submit mandatory essay on the state of proxy scalability
- Restructure essay into draft systematization of knowledge chapter
- Start development of DNS server

Autumn 2026

- Finish development of DNS server and custom ingress controller
- Implement client simulator and dummy microservices
- Select baseline paradigms
- Start developing the experiments
- Begin writing thesis background and methodology chapters

Spring 2027

- Perform the experiments
- Analyze the results
- Write article on DNS-based traffic steering
- Finish master thesis

8 Conclusion

References

- [1] *IPv6 Address Allocation and Assignment Policy*, <https://www.ripe.net/publications/docs/ripe-738/>. Accessed: May 21, 2026.
- [2] J. H. Saltzer, D. P. Reed, and D. D. Clark, “End-to-end arguments in system design,” *ACM Transactions on Computer Systems*, vol. 2, no. 4, pp. 277–288, Nov. 1984, ISSN: 0734-2071, 1557-7333. DOI: [10.1145/357401.357402](https://doi.org/10.1145/357401.357402) Accessed: May 21, 2026.
- [3] Y. Wang, C. Shou, J. Qian, and G. Liu, *XLB: A High Performance Layer-7 Load Balancer for Microservices using eBPF-based In-kernel Interposition*, 2026. DOI: [10.48550/ARXIV.2602.09473](https://doi.org/10.48550/ARXIV.2602.09473) Accessed: May 21, 2026.
- [4] M. Li, W. Lu, H. Lin, J. Wu, Y. Zhang, and G. Yan, “FlatProxy: A DPU-centric Service Mesh Architecture for Hyperscale Cloud-native Application,” vol. 14, no. 8, 2023.
- [5] K. Miziński and S. Przyłucki, “The impact of using eBPF technology on the performance of networking solutions in a Kubernetes cluster,” *Journal of Computer Sciences Institute*, vol. 35, pp. 150–158, Jun. 2025, ISSN: 2544-0764. DOI: [10.35784/jcsi.7110](https://doi.org/10.35784/jcsi.7110) Accessed: May 21, 2026.
- [6] J. Hawley, “GeoDNS—Geographically-aware, protocol-agnostic load balancing at the DNS level,”
- [7] G. C. M. Moura, J. Heidemann, R. d. O. Schmidt, and W. Hardaker, “Cache Me If You Can: Effects of DNS Time-to-Live,” in *Proceedings of the Internet Measurement Conference*, ser. IMC ’19, New York, NY, USA: Association for Computing Machinery, Oct. 2019, pp. 101–115, ISBN: 978-1-4503-6948-0. DOI: [10.1145/3355369.3355568](https://doi.org/10.1145/3355369.3355568) Accessed: May 21, 2026.
- [8] A. Shaikh, R. Tewari, and M. Agrawal, “On the effectiveness of DNS-based server selection,” in *Proceedings IEEE INFOCOM 2001. Conference on Computer Communications. Twentieth Annual Joint Conference of the IEEE Computer and Communications Society (Cat. No.01CH37213)*, vol. 3, Apr. 2001, 1801–1810 vol.3. DOI: [10.1109/INFCOM.2001.916678](https://doi.org/10.1109/INFCOM.2001.916678) Accessed: May 21, 2026.
- [9] A. Jeffery, H. Howard, and R. Mortier, “Rearchitecting Kubernetes for the Edge,” in *Proceedings of the 4th International Workshop on Edge Systems, Analytics and Networking*, ser. EdgeSys ’21, New York, NY, USA: Association for Computing Machinery, Apr. 2021, pp. 7–12, ISBN: 978-1-4503-8291-5. DOI: [10.1145/3434770.3459730](https://doi.org/10.1145/3434770.3459730) Accessed: May 21, 2026.
- [10] D. Chatterjee, Z. Tari, and A. Zomaya, “A task-based adaptive TTL approach for Web server load balancing,” in *10th IEEE Symposium on Computers and Communications (ISCC’05)*, Jun. 2005, pp. 877–884. DOI: [10.1109/ISCC.2005.19](https://doi.org/10.1109/ISCC.2005.19) Accessed: May 21, 2026.